Portfolio Optimisation with Sequential Monte Carlo

Yow Tzu Lim

Submitted in partial fulfilment of the requirements for the
MSc in Statistics

Imperial College, London

Department of Mathematics

September 1, 2014

Declaration

The work contained in this thesis is my own work unless otherwise stated.

I am the primary author of all work reported in this thesis. Advice on specific aspects of the work was provided by my supervisor, Dr. Nikolas Kantas.

Abstract

Constructing an optimal portfolio is the ultimate goal of every investment manager; but the optimality criteria may be very different to each of them. For index tracker fund manager, the main objective of portfolio management is to track and replicate the exposure of a benchmark index.

In this study, we view this as a stochastic control problem. We adopt the Bayesian view and treat the investment decision controls as random variables. The objective is to search for the sequence of control parameters that results in a portfolio that tracks a benchmark index in an optimal way. We investigate here the potential of using Sequential Monte Carlo (SMC) as the means of determining such strategy. We examine the feasibility of this approach using two examples. The first example is a toy example with the target reference set to be oscillating oscillating wave. In the second example, a real world index — the German's DAX index level is used as the target reference. In both cases, the approach looks very promising. Lastly, we demonstrate how this technique can be integrated into the Model Predictive Control (MPC) framework.

The thesis concludes with an evaluation on the work done, the extent to which the work justify the thesis hypothesis and some possible directions on how SMC can be applied to address a wider range of relevant problems on the domain of concern.

Contents

List of Figures

List of Tables

Acknowledgements

I would like to express my gratitude to all who have provided support in accomplishing this thesis. I would like to thank my supervisor, Dr. Nikolas Kantas, for the opportunity given to carry out this research, the guidance, expertise, encouragement and patience during the course of this research.

I would also like to express my appreciation to my wife Ka Ki Lai for her endless support and care in my personal life during the course of this master degree study. I would like to thank God Almighty, who has given me the strength and peace to work hard on this project. I would like to express my gratitude to all who have provided support in accomplishing this thesis.

Chapter 1

Introduction

Capital allocation is a common challenge for every investor. In the investment decision making process, investors decide how much capital is allocated to different investable assets to form a portfolio that performs better than any other possible portfolios according to some criteria. These criterion can be different among investors. Some investors may also have additional considerations such as tax considerations and legal restriction on investment assets or holding periods.

Two common objectives, which often contradicting, are the financial return and the investment risk. The Markowitz's modern portfolio theory [17] proposes a portfolio selection framework in which it is assumed that investor attempt to maximise a portfolio's return and minimise the risk (measured by the variance of the portfolio return). Based on these criteria, the set of non-dominated portfolios, commonly known as the *efficient portfolios* of a given investment period can be found. However, using the variance of the portfolio return as a risk measure has its limitation. Variance is a symmetric measure; an out-performing asset (than the expected return) is deemed to as risky as an under-performing one. Many alternative risk measurements have been proposed, e.g. Sortino ratio, Conditional Value at Risk (CVaR), etc.; see [21] for details. Two common objectives, which often contradicting, are the financial return and the investment risk. The Markowitz's modern portfolio theory [17] proposes a portfolio selection framework in which it is assumed that investor attempt to maximise a portfolio's return and minimise the risk (measured by the variance of the portfolio return). Based on these criteria, the set of non-dominated portfolios, commonly known as the *efficient portfolios* of a given investment period can be found. However, using the variance of the portfolio return as a risk measure has its limitation. Variance is a symmetric measure; an out-performing asset (than the expected return) is deemed to as risky as

an under-performing one. Many alternative risk measurements have been proposed, e.g. Sortino ratio, Conditional Value at Risk (CVaR), etc.; see [21] for details.

There are some investment fund managers who have no strong interest in maximising their portfolio's return. Instead, the main objective of portfolio management for such a fund is to simply track and replicate the exposure of a benchmark index as close as possible. These funds are attractive as it provides the investors the exposure to the market at low fees and taxes because only minimal active management is required. Passively tracking a benchmark index also makes the fund less vulnerable to the change of fund managers. The performance of these funds are often assessed in term of how well the fund tracks the benchmark index using some pre-defined metrics. These index tracking funds are the application of interest in this thesis.

1.1 Methodology

In the era before modern computing, portfolio optimisation have been explored in an analytical fashion, adopting necessary assumption as necessary. This seems rather restrictive; there are many instances where numerical methods have been used to derive an approximate or even more effective solution to the problem in question. For example, Monte Carlo methods are used to approximate integral, heuristic optimisation search techniques such as simulated annealing applied in engineering domains, etc.

In this thesis, we view a portfolio optimisation as a *stochastic* control problem. We adopt the Bayesian view and treat the controls as random variables. The objective is to find the sequence of control parameters that optimise the reward function defined in terms tracking error between the fund and the benchmark index. We investigate the potential of using Sequential Monte Carlo (SMC) as the means of determining the optimal strategies for this problem. The main reason of choosing SMC is its ability to carry out *sequential* update on the posterior distribution over time fit well with the problem in question and its success in its applications in many different domains. Of course, other heuristic search techniques are also potentially applicable.

To investigate this approach, we first used SMC to track the output of a simple deterministic reference model. This model is doubly useful. It demonstrates the concept nicely and serves as a basic model to allow us to gain further understanding on the tuneable parameters. We then considered the problem of tracking a real-world index with its constituent prices modelled as Arithmetic Brownian motion with drift. Using SMC, we search for the optimal strategy (the optimal values of the control parameters at each time point) with respect to the reward function. The reward function is defined

in terms of the tracking error and the transaction costs involved in the changes of the portfolio's holding position. Lastly, we introduce the Model Predictive Control (MPC) framework and show how to integrate this technique into the MPC framework easily.

1.2 Contributions

In this thesis, it is found that SMC has the potential to be an effective means of searching the optimal strategy for index tracking funds. The following work has been carried out:

1. exploring the potential of SMC in searching the optimal strategy for minimising the tracking error and transaction costs of an index fund.

2. exploring the sensitivity of SMC in terms of the parameter settings, the trade-off between estimation accuracy and computational efforts numerically and providing suggestions for real-world problem.

3. introducing the Model Predictive Control (MPC) framework and demonstrate how to integrate SMC technique proposed into the MPC framework.

Given the time frame of the project, we fully understand it is impossible to evaluate our approach on full scale strategy. The aim is to establish the plausibility or, at the very least, a greater understanding of the strengths and weaknesses of the above approach.

1.3 Thesis organisation

The subsequent chapters of this thesis are organised as follows:

- Chapter 2 reviews some fundamental concepts in Monte Carlo methods that are related to this thesis. It begins with a brief introduction to basic methods such as perfect Monte Carlo sampling, rejection sampling and importance sampling. It then presents Sequential Monte Carlo (SMC) methods include Sequential Importance Sampling (SIS), Sequential Importance Resampling (SIR) and various further enhancements proposed in the literature, e.g., Effective Sample Size (ESS), MCMC move, Rao-Blackwellisation, that are being used in this thesis.

- Chapter 3 briefly reviews the portfolio optimisation problem in practice. It then discusses how a portfolio problem can be transformed into a path-space parameter estimation problem. It then presents a simple experiment in which we attempt to use SMC to track a reference signal generated by the model.

- Chapter 4 details the application of the algorithm to track the DAX index. It presents several experiments that have been carried out to demonstrate how the index can be tracked with full replication as well as partial replication. Lastly, it introduces Model Predictive Control (MPC) framework and demonstrates how this technique can be integrated into the MPC framework easily.

- Chapter 5 concludes the thesis by evaluating the contributions of the thesis and discusses some potential work for the future.

Chapter 2

Monte Carlo methods

Monte Carlo Methods have achieved significant success in its application to various domains in the last few decades. This chapter reviews some fundamental concepts in Monte Carlo methods that are related to this thesis. It begins with a summary on the main concept of Bayesian inference. It then discusses some basic Monte Carlo methods such as perfect Monte Carlo sampling, rejection sampling and importance sampling. Lastly, it introduces Sequential Monte Carlo (SMC), along with various enhancement made to the techniques, used in this thesis.

2.1 Bayesian inference

In Bayesian inference, each unknown parameter in the model is assumed to be random variable and is associated with a prior distribution that characterises the initial belief. The inference process is merely about updating the belief with new observable evidence in a systematic fashion using Bayes theorem.

Formally, let $\mathcal{M}$ be the model used, θ be the set of parameters of the model, $p\left(\theta \mid \mathcal{M}\right)$ be the prior distribution (initial belief) and $p(x \mid \theta, \mathcal{M})$ be the likelihood (probability of observing an observation x given the model) then posterior distribution (updated belief) is given as follows:

$$p(\theta \mid x, \mathcal{M}) = \frac{p(x \mid \theta, \mathcal{M}) \ p(\theta \mid \mathcal{M})}{p(x \mid \mathcal{M})}$$

$$\propto p(x \mid \theta, \mathcal{M}) \ p(\theta \mid \mathcal{M}) \qquad (2.1)$$

$$\text{posterior} \propto \text{likelihood} \times \text{prior} \qquad (2.2)$$

This problem formulation is elegant, but there remains some subtle issues in practice. One particular issue is the calculation of the normalisation constant $p(x \mid \mathcal{M})$ in

(2.1), which demands us to be able to carry out the following integral analytically:

$$p(x \mid \mathcal{M}) = \int p(x \mid \theta, \mathcal{M}) p(\theta \mid \mathcal{M}) \ d\theta \tag{2.3}$$

This is often infeasible. A often way to circumvent this requirement is by making use of conjugate prior that yields posterior distributions from the same family in an analytical fashion. Moreover, the need of calculating integral that does not possess analytic solution also arises in the marginalisation process of nuisance parameters, calculating expectation of a function, etc.

2.2 Perfect Monte Carlo

Instead of a closed form solution, the Method Carlo methods offer a numerical solution in estimating the integral using simple sampling techniques. Consider the calculation the expectation of a function, I of the following form:

$$I = \mathbb{E}_p[f(x)] = \int f(x)p(x) \ dx \tag{2.4}$$

Assuming we are able to sample N independent and identically distributed (i.i.d.) samples of x from $p(\cdot)$, denote these as $\{x^{(i)}\}$ where $i \in \{1 \dots N\}$, a Monte Carlo estimate of I using the the point masses of the samples is:

$$\hat{I} = \frac{1}{N} \sum_{i=1}^{N} f(x^{(i)}) \tag{2.5}$$

This approximation can be viewed as discretisation of the continuous distribution with *random* grid points. This estimate is unbiased and converge almost surely to I as $N \to \infty$ by the Law of Large number [20].

Moreover, if the variance of $f(\cdot)$ is bounded ($\sigma_f^2 < \infty$), as $N \to \infty$, then the following central limit theorem holds:

$$\sqrt{N}(\hat{I} - I) \Longrightarrow N(0, \sigma_f^2) \tag{2.6}$$

where $\Longrightarrow$ denotes convergence in distribution [7]. The key point to note here is this convergence rate of $\frac{1}{\sqrt{N}}$ is independent of the dimensions of x. This is in contrast with any deterministic method that has a rate that decreases as the integral dimension increases [20]. This is the main advantage of Monte Carlo integration.

2.3 Rejection sampling

However, it is not always possible to sample directly from the distribution $p(\cdot)$. Suppose we can find an instrumental distribution (a.k.a. proposal distribution), $q(\cdot)$, that is easy to sample from and has the property such that $cq(x)$ dominates $p(x)$ for all x, i.e., $cq(x) \geq p(x) \geq 0$ for all x, then to get a random sample from $p(\cdot)$, we can first sample from $q(\cdot)$ instead and accept the sample with acceptance probability $\alpha(x) = \dfrac{p(x)}{cq(x)}$. If the sample is rejected, the process is repeated until success. This rejection algorithm is summarised in Algorithm 1[1].

Algorithm 1 Rejection Sampling

1: **function** REJECTIONSAMPLING(N)
2:     $\mathcal{X} = \{\ \}$.
3:     **repeat**
4:         sample $x \sim q(\cdot)$.
5:         sample $u \sim \mathcal{U}(0,1)$.
6:         **if** $u \leq \dfrac{p(x)}{cq(x)}$ **then**.
7:             $\mathcal{X} \leftarrow \mathcal{X} \cup \{x\}$.
8:         **end if**
9:     **until** len($\mathcal{X}$)=N.
10:     return $\mathcal{X}$.
11: **end function**

Looking at the acceptance ratio formula, it is not difficult to see that the optimal instrumental distribution, q^*, is the one that minimises the space bounded by $cq(x)$ subject to the constraint that it still dominates the target density $p(x)$. As the dimension of x increases, this algorithm becomes very inefficient because the acceptance ratio which is essentially defined as the ratio of two spaces tends towards zero. Therefore, many generated examples would be rejected.

2.4 Importance sampling

Instead of making a binary accept-reject decision on each sample, the key concept in important sampling is assign weighting to each sample (obtained from the instrumental distribution, $q(\cdot)$) based on how well the sample resembles the target distribution, $p(\cdot)$. More formally, assume we have an instrumental distribution, $q(\cdot)$ that has support that

---

[1]The set notation is used here, despite it is actually a collection that allows duplicates. Therefore, the $\cup$ operation should be viewed as a simple join of two collections.

includes $p(\cdot)$, we can re-write (2.4) as:

$$
\begin{aligned}
I &= \int f(x)\frac{p(x)}{q(x)}q(x)\ dx \\
&= \int f(x)w(x)q(x)\ dx \\
&= \mathbb{E}_q[f(x)w(x)]
\end{aligned}
\tag{2.7}
$$

where $w(x)$ is commonly referred as the importance weight. This reformulation leads to the following Monte Carlo estimate of I:

$$
\begin{aligned}
\hat{I} &= \frac{\frac{1}{N}\sum_{i=1}^{N}\tilde{w}(x^{(i)})f(x^{(i)})}{\frac{1}{N}\sum_{j=1}^{N}\tilde{w}(x^{(j)})} \\
&= \sum_{i=1}^{N}\frac{\tilde{w}(x^{(i)})}{\sum_{j=1}^{N}\tilde{w}(x^{(j)})}f(x^{(i)}) \\
&= \sum_{i=1}^{N}\hat{w}(x^{(i)})f(x^{(i)})
\end{aligned}
\tag{2.8}
$$

where $\tilde{w}(x^{(i)}) = \frac{p(x^{(i)})}{q(x^{(i)})}$ and $\hat{w}(x^{(i)}) = \frac{\tilde{w}(x^{(i)})}{\sum_{j=1}^{N}\tilde{w}(x^{(j)})}$ are referred to as unnormalised and normalised importance weight respectively [4]. This estimate is biased as it consists of the ratio of two estimates, yet it is still asymptotically consistent.

To obtain samples from the target distribution, $p(\cdot)$, an additional resampling step can be introduced. In the first step, we draw a set of samples $\{\tilde{x}^{(i)}\}$ from the instrumental distribution and compute their associated normalised importance weights, $\hat{w}(\tilde{x}^{(i)})$. In the resampling step, we draw the final sample set, $\{x^{(i)}\}$ from this intermediate set of samples by taking into account the importance weights. This importance sampling algorithm is summarised in Algorithm 2

2.5 Sequential Monte Carlo

To motivate why Sequential Monte Carlo is useful, consider a target distribution of interest $p(x_{0:n})$, and for simplicity, assuming we can sample directly from the distribution $p(x_{0:n})$, the minimal computational complexity of the sampling scheme would be at least linear in n. Sequential Monte Carlo (SMC) provides a way to obtain samples of x for each sequential time step in a *fixed* amount of computational time in Hidden Markov Models (HMMs). We shall begin with a brief introduction on HMMs that is crucial to understand SMC in the next section. Refer [4] for further details of inference techniques for HMMs in general.

Algorithm 2 Importance Sampling

1: **function** IMPORTANCESAMPLING(N)
2:      $\tilde{\mathcal{X}} = \{\ \}$
3:      **repeat**
4:          sample $\tilde{x} \sim q(\cdot)$.
5:          $\tilde{\mathcal{X}} \leftarrow \tilde{X} \cup \{\tilde{x}\}$.
6:      **until** len($\tilde{X}$)=N.
7:      calculate importance weights, $\hat{w}(\tilde{x}^{(i)}) = \dfrac{\tilde{w}(\tilde{x}^{(i)})}{\sum_{j=1}^{N} \tilde{w}(\tilde{x}^{(j)})}$.
8:      $\mathcal{X} = \{\ \}$.
9:      **repeat**
10:         sample x from $\tilde{\mathcal{X}}$ according to the importance weights, $\hat{w}$.
11:         $\mathcal{X} \leftarrow \mathcal{X} \cup \{x\}$.
12:     **until** len($\mathcal{X}$)=N.
13:     return $\mathcal{X}$.
14: **end function**

2.5.1   Hidden Markov Models

HMMs can be seen as a class of models that consist of two related processes: an underlying Markov process, X_t, which is the target process of interest, and a observable process, Y_t, which its state can be measured and therefore provides some information about X_t. Moreover, it is assumed that these two processes have conditional independence properties as shown using the graphical model representation in Figure 2.1. These properties can be summarised as follows:

$$p(x_t \mid x_{0:t-1}) = f(x_t \mid x_{t-1})$$
$$p(y_t \mid x_{0:t}, y_{0:t-1}) = g(y_t \mid x_t) \tag{2.9}$$

where $f(x_t \mid x_{t-1})$ is the Markov transition density and $g(y_t \mid x_t)$ is the conditionally independent likelihood.

Figure 2.1: Hidden Markov Models

This class of models are designed to model evolving systems that output some observable events over time, which depends on the internal state of the system in a on-line manner. This imposes an implicit requirement on computation perspective: the estimate calculation cost should remain constant over time, i.e., the calculation cost does not increase with the increasing number of states.

Arguably, the most common inference problem in HMMs is the smoothing distribution, $p(x_{0:t} \mid y_{0:t})$, that is estimating the states $x_{0:t}$ based on the sequence of observations up to time t, $y_{0:t}$. Using Bayes rules, we can write the density of the distribution of interest in the recursion form as follows:

$$
\begin{aligned}
p(x_{0:t} \mid y_{0:t}) &= \frac{p(x_{0:t}, y_{0:t})}{p(y_{0:t})} \\
&= \frac{p(x_{0:t-1}, y_{0:t-1}) f(x_t \mid x_{t-1}) g(y_t \mid x_t)}{p(y_{0:t})} \\
&= p(x_{0:t-1} \mid y_{0:t-1}) \frac{f(x_t \mid x_{t-1}) g(y_t \mid x_t)}{p(y_t \mid y_{0:t-1})}
\end{aligned} \tag{2.10}
$$

This recursion is often re-written into two separate steps: the prediction step (the estimation of distribution of all states up to time t given only states up to time $t-1$) and the update step (the correction of the predicted distribution taking into account the new observation) as follows:

$$
\begin{aligned}
p(x_{0:t} \mid y_{0:t-1}) &= p(x_{0:t-1} \mid y_{0:t-1}) f(x_t \mid x_{t-1}) \\
p(x_{0:t} \mid y_{0:t}) &= \frac{p(x_{0:t} \mid y_{0:t-1}) g(y_t \mid x_t)}{p(y_t \mid y_{0:t-1})}
\end{aligned} \tag{2.11}
$$

Moreover, any smoothing distribution $p(x_{j:k} \mid y_{1:t})$ where $(0 \leq j \leq k \leq t)$ can be obtained by integrating out x_i that are not interested in as follows:

$$
p(x_{j:k} \mid y_{0:t}) = \int p(x_{0:t} \mid y_{0:t}) dx_{0:j-1, k+1:t} \tag{2.12}
$$

A common smoothing distribution of interest is the marginal distribution at time t, $p(x_t \mid y_{0:t})$, which is also referred to as the filtering distribution.

Another distribution of interest is the prediction distribution, that is the estimation of the distribution of any unseen *future* states using only all observations up to current time. If we let $j = 0$ and $k \geq t$ in (2.12), we obtain the following equation:

$$
p(x_{0:k} \mid y_{0:t}) = p(x_{0:t} \mid y_{0:t}) \prod_{i=t+1}^{k} f(x_i \mid x_{i-1}) \tag{2.13}
$$

Therefore, any prediction density can be obtained by simply integrating out the variables of not interest from the above equation.

While these expressions *appear* to be very simple, the distribution estimation problem is in fact far from being resolved in practice. The integrals involved in the above equations are often intractable because they require the evaluation of complex high-dimensional integrals[1].

2.5.2 Sequential Important Sampling (SIS)

Sequential Monte Carlo provides a systematic way to approximate the solution to the distribution estimation problem. Let assume that it is possible to decompose the selected proposal distribution into recursion form as follows:

$$q_{0:t}(x_{0:t} \mid y_{0:t}) = q_{0:t-1}(x_{0:t-1} \mid y_{0:t-1})q_t(x_t \mid x_{0:t-1}, y_{0:t})$$

$$= q_0(x_0 \mid y_0) \prod_{i=1}^{t} q_i(x_i \mid x_{0:i-1}, y_{0:i}) \tag{2.14}$$

then it is possible to obtain sample $x_{0:t}$ by first sampling $x_0 \sim q_0(\cdot)$ at time 0 and then $x_i \sim q_i(x_i \mid x_{0:i-1}, y_{0:i})$ for all time i from 1 to t. The corresponding weight associated to each sample $x_{0:t}$ can also be decomposed into recursion form as follows:

$$\tilde{w}_t = \frac{p_{0:t}(x_{0:t} \mid y_{0:t})}{q_{0:t}(x_{0:t} \mid y_{0:t})}$$

$$\propto \tilde{w}_{t-1} \frac{f_t(x_t \mid x_{t-1})g_t(y_y \mid x_t)}{q_t(x_t \mid x_{0:t-1}, y_{0:t})} \tag{2.15}$$

$$\propto \tilde{w}_0 \prod_{i=1}^{t} \alpha_i(x_i) \tag{2.16}$$

where $\alpha_t(x_t)$ is often referred to as incremental importance weight function. This is the key concept in SIS.

Using these weighted samples $\{(x_{0:t}^{(i)}, \tilde{w}_t^{(i)})\}_{1 \leq i \leq N}$, it is possible to estimate any function f defined on the space using the self normalised importance sampling estimator in the same way as (2.8) as follows:

$$\hat{f}(x_{0:t}) = \sum_{i=1}^{N} \frac{\tilde{w}_t^{(i)}}{\sum_{j=1}^{N} \tilde{w}_t^{(j)}} f(x_{0:t}^{(i)})$$

$$= \sum_{i=1}^{N} \hat{w}^{(i)} f(x_{0:t}^{(i)}) \tag{2.17}$$

---

[1]Except under very special setting, in which the marginal filtering density $p(x_t \mid y_{0:t})$, the prediction density $p(x_t \mid y_{0:t-1})$ and the recursive likelihood $P(y_t \mid y_{0:t-1})$ are all Gaussian densities, then their means and variances can be computed analytically using Kalman Filter as shown in Appendix A.

A particular important instance of SMC is obtained when the prior distribution is used as the importance distribution:

$$q_{0:t}(x_{0:t} \mid y_{0:t}) = p(x_{0:t})$$
$$= p_0(x_0) \prod_{i=1}^{t} p_i(x_i \mid x_{0:i-1}) \tag{2.18}$$

This greatly simplifies importance weight to follows:

$$\tilde{w}_t \propto \tilde{w}_{t-1} g(y_t \mid x_t) \tag{2.19}$$

The SIS algorithm is summarised in Algorithm 3.

Algorithm 3 Sequential Importance Sampling

1: **function** SEQUENTIALIMPORTANCESAMPLING(N, T)
2:     Set $t \leftarrow 0$.
3:     For $i \in 1, \ldots, N$, sample $x_0^{(i)} \sim q(x_0^{(i)} \mid y_0^{(i)})$.
4:     For $i \in 1, \ldots, N$, calculate the unnormalised importance weights:

$$\tilde{w}_0^{(i)} = \frac{f(x_0^{(i)}) g_0(y_0 \mid x_0^{(i)})}{q_0(x_0^{(i)} \mid y_0)}$$

5:     For $i \in 1, \ldots, N$, normalize the importance weights:

$$\hat{w}_0^{(i)} = \frac{\tilde{w}_0^{(i)}}{\sum_{i=1}^{N} \tilde{w}_0^{(i)}}$$

6:     Set $t \leftarrow t + 1$.
7:     **while** $t \leq T$ **do**
8:         For $i \in 1, \ldots, N$, sample $x_t^{(i)} \sim q(x_t^{(i)} \mid y_{t-1}, x_{t-1}^{(i)})$.
9:         For $i \in 1, \ldots, N$, set $x_{0:t}^{(i)} \leftarrow (x_{0:t-1}^{(i)}, x_t^{(i)})$.
10:        For $i \in 1, \ldots, N$, calculate the unnormalised importance weights:

$$\tilde{w}_t^{(i)} = w_{t-1}^{(i)} \frac{f_t(x_t^{(i)} \mid x_{t-1}^{(i)}) g_t(y_t \mid x_t^{(i)})}{q_t(x_t^{(i)} \mid x_{0:t-1}^{(i)}, y_{0:t})}$$

11:        For $i \in 1, \ldots, N$, normalize the importance weights:

$$\hat{w}_t^{(i)} = \frac{\tilde{w}_t^{(i)}}{\sum_{i=1}^{N} \tilde{w}_t^{(i)}}$$

12:    **end while**
13: **end function**

2.5.3 Optimal proposal distribution

While SIS is attractive, it is nothing but a specialised version of importance sampling introduced earlier in Section 2.4. As the state space increases with the number of time step t, direct importance sampling on a state space that is increasing in size is not very efficient. The weights of the samples will degenerate over time, in the sense that the weights start to concentrate only on a small number of samples. Consequently, many samples will have negligible weights and do not contribute much in the estimating the expectation. See [10] for a step-by-step illustration. The weight degeneracy issue cause the quality of the estimate degrades over time.

To alleviate this issue, looking at (2.15), it is obvious that the variance of importance weights can be minimised by using a proposal distribution of the following form:

$$p_t(x_t \mid x_{t-1}, y_t) \propto f_t(x_t \mid x_{t-1})g_t(y_t \mid x_t) \qquad (2.20)$$

This is often referred to as the optimal proposal distribution.

In general, it is not always possible to sample from this optimal proposal distribution. Yet, the knowledge of its form can be helpful in designing a reasonable good proposal distribution, which one can sample from. Better proposal *reduces* the amount of variance introduced, but it *does not eliminate* the weight degeneracy problem.

2.5.4 Sequential Importance Resampling (SIR)

The variance in importance weights accumulates over iterations. This suggests a possible solution is to "reset" the weights associated to the samples somehow during the iterations [10]. Sequential Importance Resampling (SIR) introduces an additional resampling step to SIS step in a similar fashion as discussed in Section 2.4. After resampling, the weight of each sample is reset to be equal, i.e., $\frac{1}{N}$. This resampling step however introduces extra variance to the estimators.

A simple direct implementation of resampling step is to select the sample from the intermediate stage according to a Multinomial distribution with the success probability parameter set to the vector of normalised weights, $\hat{w}(x^{(i)})$, i.e., the chance of a sample point being replicated is proportional to its weight.

There are many other resampling schemes have been proposed in the literature. For example, stratified resampling [14] as the name suggested splitting the samples into strata to ensure the good coverage on the resulting sample set, residual resampling [15] which is able to decrease the variance of the weights due to resampling, etc. See [6] for further details on the comparison of these sampling schemes. The SIR algorithm is now summarised in Algorithm 4.

2.5.5 Effective sample size (ESS)

Resampling step induces additional Monte Carlo variance to the weights. Yet, this step is necessary to avoid accumulation of estimation variance onto the weights over time and therefore results in a more accurate estimate.

To trade-off these two competing requirements, one possible way is to monitor the effective sample size (ESS) of the sample set which provides a measure on the quality of the weighted samples. The ESS value can be estimated as follows:

$$\text{ESS} \approx \frac{1}{E[w^2]} \approx \frac{\left(\sum_{i=0}^{N} w_i\right)^2}{\sum_{i=0}^{N} w_i^2} \tag{2.21}$$

A common way to integrate this idea is to trigger the resampling step only if the ESS value at time step t fall below certain threshold at time t, say $\frac{N}{2}$. See [10] for further details on ESS.

Algorithm 4 Sequential Importance Resampling

1: **function** SEQUENTIALIMPORTANCERESAMPLING(N, T)

2:     Use Steps 2-11 of Algorithm 3

3:     **Resample:** For $i \in 1, \dots, N$, resample $x_{0:t}^{(i)} \sim \dfrac{\sum_{i=1}^{N} \hat{w}_t^{(i)} \delta_{x_{0:t}^{(i)}}}{\sum_{j=1}^{N} \hat{w}_t^{(j)}}$

4: **end function**

2.5.6 Resample-Move Algorithm

However, resampling is not a silver bullet for sampling impoverishment. Essentially, resampling provides a mechanism to eliminate low weight samples to give way to replicate *copies* of high weight samples. This allows all samples to participate and contribute to the distribution estimation. This is obvious beneficial for the case of estimating filtering distribution and predictive distribution. Over time, this replication results in decrease in the number of distinct samples for previous time steps. Eventually, many samples will have the share the same sample trajectory. This is a fundamental weakness of SMC, in which the sample path history is never re-written. This lose of diversity in the sample set will have a negative impact when it comes to any smoothing distribution estimation.

To counteract this sample impoverishment, Resample-Move Algorithm [3, 8] is proposed to introduce some perturbation to the samples (so to diversify them) without changing the distribution they represent. In the original paper, this is accomplished by adding a MCMC "move" step to each sample after the resampling step with a Markov

Kernel, K that is invariant to the target distribution, $p(\cdot)$. This Resample-Move algorithm is summarised in Algorithm 5.

This does not entirely solve the smoothing distribution estimation problem. To apply Markov Kernels with invariant distribution corresponding to the smoothing distribution, the space that Markov kernel is defined has to increase at each iteration. This implies the computation time increases linearly with time. Moreover, fast mixing high dimension Markov kernel in itself is not easy to design [10].

To trade-off between the two requirements, one could adopt a sliding windows approach, in which the MCMC Kernel is used to diversify only the samples of the previous n time step at each iteration. Adding this sliding window approach into the standard SIR algorithm will make each iteration has an additional *fixed* computational cost.

Algorithm 5 Resample-Move Algorithm

1: **function** RESAMPLEMOVEALGORITHM(N, T)
2:      Use Steps 2-11 of Algorithm 3
3:      **Resample:** For $i \in 1, \dots, N$, resample $x_{0:t}^{(i)} \sim \dfrac{\sum_{i=1}^{N} \hat{w}_t^{(i)} \delta_{x_{0:t}^{(i)}}}{\sum_{j=1}^{N} \hat{w}_t^{(j)}}$
4:      **Move:** For $i \in 1, \dots, N$, sample $x_{0:t}^{(i)} \sim K_t(\cdot)$, where K_t is p_t-invariant.
5: **end function**

2.5.7   Rao-Blackwellised SMC

In practice, many models may not be entirely non-linear or non-Gaussian. Some states may be linear and Gaussian, conditional upon other states. Instead of naively modelling all state spaces using SMC, a better approach would be carrying out analytical computation as possible on the linear Gaussian states and use SMC model only the non-linear states. This marginalisation setting will yield estimate that has smaller variance.

To illustrate this, let consider the a simple conditional linear Gaussian Model as follows:

$$X_t^L = A_t(X_t^N)X_{t-1}^L + B_t(X_t^N)W_t + F_t(X_t^N) \tag{2.22}$$

$$Y_t = C_t(X_t^N)X_t^L + D_t(X_t^N)V_t + G_t(X_t^N) \tag{2.23}$$

where $\{X_t^N\}_{t \geq 0}$ is a non-linear Markov process, A_t, B_t, C_t, D_t, F_t, G_t are appropriate matrix/vector of X_t^N and $\{W_t\}_{t \geq 0}$ and $\{V_t\}_{t \geq 0}$ are independent sequences of standard Gaussian random variable, i.e., $W_t, V_t \sim \mathcal{N}(0, I)$. In such a case, the transition density and likelihood of this model are Gaussian distributions with centre lied at a point of

a linear combination of the known x_t^N that can be obtained by using classical Kalman Filtering recursions as shown in Appendix A.

The discussion here has been focused on conditional linear Gaussian model. The discrete state space HMMs is another important class of HMMs in which HMM forward algorithm [18] can be used to marginalise the internal states. See [4] for further details.

2.6 Conclusion

This chapter presents a review of Monte Carlo methods, with a particular focus on Sequential Monte Carlo (SMC), along with various extensions proposed in the literature, used in this thesis for portfolio optimisation.

It is worth to have a remark here that the SMC are not only applicable to sequential filtering problem. For example, it has been established that it is possible to use SMC within MCMC framework (pMCMC, where p stands for particle) [2] to build high dimensional proposal distribution to improve standard MCMC. In the next chapter, we will show how SMC can be used as a maximiser to search for a optimal strategy for the optimal portfolio, given a multiplicative reward function.

Chapter 3

Portfolio optimisation

Optimal portfolio is the ultimate goal of every investment manager; but the optimality criteria can be very different to each of them. In the infamous Markowitz's modern portfolio theory [17], it is assumed that investor attempt to maximise the return and minimise the risk of his portfolio. For index tracker fund manager, the main objective of portfolio management is to track and replicate the exposure of a benchmark index. The lack of active management generally makes the fund less vulnerable to change of management and has the advantages of lower fees and taxes. It is the latter the focus of the thesis lies.

In this chapter, we introduces the index tracking fund and summarise some causing factors of the tracking error between the fund and its benchmark index. We then present how the portfolio optimisation problem in terms of minimising the tracking error can be formulated as a stochastic control problem. Lastly, we show this stochastic control problem can be turned into a path-space parameter estimation problem, in which Rao-Blackwellised SMC can be used to estimate these parameters efficiently.

3.1 Index tracking fund

An financial index can be thought of a summary statistics of a market, typically formed as weighted average of the prices of some financial instruments. For example, FTSE 100 is an index that attempts to represents the performance of the 100 biggest companies in UK. This index is however not an investable financial product one can invest directly.

To allows investors to take exposure to the markets (as represented by the indices), index tracker funds are introduced. These funds generally follow very strict and transparent rules with the main objective to track their benchmark indices as close as possible. Investors of these funds therefore are only exposed to the market risks of

the chosen indices, with minimal exposure to the risk associated with the investment process of the traditional active fund management.

3.1.1 Index replication

To track an index, the simplest way is to replicate the index by investing all the components of the benchmark index. However, this can be costly and difficult to manage, considering some of the indices may have huge number of components, e.g., MSCI World index that consists of ≈ 1600 components from many different countries. Instead, one could partial replicate the index by sampling some of the components that are most representative. This can mean those components with larger weights, more volatiles or less correlated in returns (assuming these returns average out the gain/loss among them). This partial replication could save transaction cost, at the same time, introducing some tracking errors to the funds.

Instead of physically replicating the index by investing in the components, the portfolio managers also have the option to enter a swap agreement with a counter-party (typically investment bank) that provides the exact return of the stock market or commodity it's tracking. Essentially, this transfers the market risk to the counter-party, at the same introducing the counter-party default risk. This technique is known synthetic replication.

3.1.2 Causing factors of tracking error

To evaluate the performance of index tracking fund, different metrics have been introduced to quantify the mismatch between the performance of the fund and its benchmark index. For example, tracking difference is the sum of absolute difference in returns between of the fund and its benchmark index. Here, we adopt the tracking error as our metric, which is defined to be the standard deviation of the absolute difference of the returns of the fund and the benchmark index defined in [11] as follows:

$$\epsilon = \sqrt{\text{Var}[r_p - r_b]} \tag{3.1}$$

where r_p is the return of the portfolio and r_b is the return of the benchmark index.

There are many factors that can cause tracking errors, some of which are summarised as follows:

1. benchmark index rebalance — the benchmark index is re-weighting its constituents periodically to reflex the changes on the market based on its methodology. To track the index, the fund has to adjust its portfolio accordingly. This

will incur some transaction costs. During the index rebalance period, cash drag may also happen between the liquidation of the constituents that have weights reduced/dropped and the addition of the constituents that have weights increased/added. This cash is essentially not participating in the market and therefore it does not reflex the changes on the benchmark index.

2. replication and sampling techniques — funds may choose to replicate the benchmark index by selecting a subset of the constituents (often the ones with larger weights and more liquid) in an effort to minimise the transaction costs. This exclusion of the smaller, less liquid constituents may introduce another source of tracking error, especially under stressed market.

3. different assumption on dividend reinvestment and taxation — the benchmark index calculation often assumes immediate reinvestment of dividends on ex-dividend dates but the fund may only able to reinvest the dividend after receiving it. The tax treatment that is applied on the dividends may also be different.

4. total expense ratio — there is an additional expense charged to the fund on daily basis to cover the management cost.

This list is by no mean exclusive. See [11] for further details.

3.2 Methodology

Instead of using a traditional deterministic state space models, we choose to use a *stochastic* modelling approach to carry out the portfolio optimisation. In particular, we focus here the the problem in minimising the tracking error between a portfolio and its benchmark index. Our aim is to determine what investment actions (buy or sell) a portfolio manager needs to do for each index component on a daily basis across the investment horizon to track the benchmark index well.

We adopt here conditional Gaussian linear state space model as follows:

$$
\begin{aligned}
X_t &= A_t(U_t)X_{t-1} + B_t(U_t)W_t + F_t(U_t) \\
Y_t &= C_t(U_t)X_t + D_t(U_t)V_t + G_t(U_t)
\end{aligned}
\tag{3.2}
$$

where $\{U_t\}_{t \geq 0}$ is a deterministic control input sequence that is used regulate the hidden states, $A_t, B_t, C_t, D_t, F_t, G_t$ are appropriate matrix/vector functions of U_n and $\{W_t\}_{t \geq 0}$ and $\{V_t\}_{t \geq 0}$ are independent sequences of standard Gaussian random variable, i.e., $W_t, V_t \sim \mathcal{N}(0, I)$. The transition density and likelihood of this model are Gaussian

distributions with centre lied at a point of a linear combination of the known conditional control parameters, u_t of the following form:

$$p_t(u_t \mid u_{t-1}) = \text{(any given form)}$$
$$f_t(x_t \mid x_{t-1}, u_t) = \mathcal{N}(A_t(u_t)x_{t-1} + F_t(u_t), B_t(u_t)B_t(u_t)^T)$$
$$g_t(y_t \mid x_t, u_t) = \mathcal{N}(C_t(u_t)x_t + G_t(u_t), D_t(u_t)D_t(u_t)^T) \tag{3.3}$$

Based on this model, the optimisation objective is to search for a sequence of controls $u_{1:t}$ that would result in a sequence of observations $y_{1:t}$ that tracks the reference signal $y_{1:T}^{ref}$ as close as possible. This problem is often known as stochastic regulation problem. We adopt the finite horizon multiplicative reward function proposed in [13] here:

$$J(u_{1:T}, y_{1:T}^{ref}, x_0) = \mathbb{E}_{x_0}\left[\exp\left(-\frac{1}{2}\sum_{t=1}^{T}\left(||y_t^{ref} - C_t(u_t)x_t - G_t(u_t)||_{Q_t(u_t))^{-1}}^2 + ||u_t - u_{t-1}||_{L_t}^2\right)\right)\right] \tag{3.4}$$

where $Q(u_t) = D_t(u_t)D_t(u_t)^T$ and L_t are assumed to known. The expectation here is taken with respect to the whole path of the Markov Chain $\{X_t\}_{t\geq0}$ with the starting point $X_0 = x_0$, i.e., $E_{x_0}[\phi(X_{1:t})] = \int \phi(X_{1:T})\prod f_t(x_t \mid x_{t-1})\ dx_{1:t}$.

The corresponding optimal open loop policy is:

$$u_{1:T}^* = \arg\max_{u_{1:T}} J(u_{1:T}; y_{1:T}^{ref}; x_0) \tag{3.5}$$

with the assumption that the maximum is attainable. This reward function is closely related to the risk sensitive control discussed in [24], with slightly difference. This reward function used here is absence of an explicitly risk sensitive constant and it also assumes the presence of $y_{1:T}^{ref}$ as the reference target instead of the observable states. Nevertheless, a risk sensitivity constant could still be introduced through D_n and L_n if necessary.

3.2.1 Problem formulation

We assume here that the sequence of controls $\{U_t\}_{t\geq0}$ is a Markov process with transition distribution $p(u_t \mid u_{t-1})$. The objective is to compute the marginal posterior distribution density:

$$\pi_t(u_{1:t}) = p(u_{1:t} \mid y_{1:t}^{ref}) \tag{3.6}$$

by standard marginalisation the distribution $p(x_{1:t}, u_{1:t} \mid y_{1:t})$ which in itself satisfies the following recursive equation:

$$p(x_{1:t}, u_{1:t} \mid y_{1:t}) = p(x_{1:t-1}, u_{1:t-1} \mid y_{1:t-1})\frac{p(y_t \mid x_t, u_t)p(x_t, u_t \mid x_{t-1}, u_{t-1})}{p(y_t \mid y_{1:t-1})} \tag{3.7}$$

One straight-forward way to approximate the recursion is directly apply SMC algorithm by consider the hidden states is a sequence of paired states $\{(U_t, X_t)\}_{t \geq 0}$. A better approach would be using the Rao-Blackwellised SMC mentioned in Section 2.5.7, attempts to do as much computation analytically as possible, by splitting the states into the linear Gaussian states and the non-linear states. Then, the Kalman Filter can be used to model and track the linear Gaussian states and the SMC is only required to model and track the non-linear states. This marginalisation setting often yields estimates with smaller variance.

To use Rao-Blackwellised SMC on this model, we consider the following factorisation on the full posterior distribution as follows:

$$p(u_{1:t}, x_{1:t} \mid y_{1:t}) = p(x_{1:t} \mid u_{1:t}, y_{1:t})p(u_{1:t} \mid y_{1:t}) \tag{3.8}$$

Looking at the right hand side of the equation, the first term is obviously Gaussian and can therefore can be estimated optimally using Kalman Filter recursions described in Appendix A. For the second term, we can rewrite it into the following recursion form:

$$\begin{aligned} p(u_{1:t} \mid y_{1:t}) &\propto p(y_t \mid u_{1:t}, y_{1:t-1})p(u_{1:t} \mid y_{1:t-1}) \\ &= p(y_t \mid u_{1:t}, y_{1:t-1})p(u_t \mid u_{1:t-1}, y_{1:t-1})p(u_{1:t-1} \mid y_{1:t-1}) \\ &= p(u_{1:t-1} \mid y_{1:t-1})p(y_t \mid u_{1:t}, y_{1:t-1})p(u_t \mid u_{1:t-1}, y_{1:t-1}) \end{aligned} \tag{3.9}$$

Assuming that it is also possible to decompose the selected importance density into recursion form as follows:

$$q_{1:t}(u_{1:t} \mid y_{1:t}) = q_{1:t-1}(u_{1:t-1} \mid y_{1:t-1})q_t(u_t \mid u_{1:t-1}, y_{1:t}) \tag{3.10}$$

then the associated unnormalised weight of each sample can be derived in a similar fashion as given in Section 2.5.2 into the recursion form as follows:

$$\tilde{w}_t = \tilde{w}_{t-1} \frac{p_t(u_t \mid u_{t-1})p_t(y_t \mid u_{1:t}, y_{1:t-1})}{q_t(u_t \mid u_{1:t-1}, y_{1:t})} \tag{3.11}$$

Compare to (2.15), the main difference here is that the weights now depends on the whole path space from time step 1 to t.

3.2.2 Extensions

As discussed in Section 2.5.5, the resampling step introduces variance to the the weights, but it is necessary to avoid accumulation of estimation variance onto the weights over time. Effective sample size (ESS) can be used to measure the quality of the weight and

trigger the resampling step only if necessary, e.g., when the ESS value at time step t fall below certain threshold at time t, say $\frac{N}{2}$.

To add diversity in the population of particles, the Resample-Move step introduced in Section 2.5.6 can be included. This step perturbs the samples yet leave the distribution the samples represent unchanged using a MCMC kernel that is invariant in density. With this optional step in place, the whole algorithm is summarised in Algorithm 6.

3.2.3 MAP estimation for the control sequence

Based on the model defined in (3.2), the conditional likelihood density of the model is:

$$g_t(y_t^{ref} \mid x_t, u_t) \propto \exp\left(-\frac{1}{2}||y_t^{ref} - C_t(u_t)x_t - G_t(u_t)||^2_{(D_t(u_t)D_t(u_t)^T)^{-1}}\right) \qquad (3.13)$$

and using the standard Kalman recursion discussed in Appendix A, the posterior distribution $p(y_t^{ref} \mid y_{t-1}^{ref}, u_{1:t})$ is given as follows:

$$p(y_t^{ref} \mid y_{t-1}^{ref}, u_{1:t}) \propto \exp\left(-\frac{1}{2}||y_t^{ref} - y_{t|t-1}(u_{1:t})||^2_{S_t(u_{1:t})^{-1}}\right) \qquad (3.14)$$

where $y_{t|t-1}$ and S_t are the mean and covariance matrix of the measurement at time t obtained in the Kalman Filter recursions.

By assuming the Markov transition density for u_t to be:

$$p(u_t \mid u_{t-1}) \propto \exp\left(-\frac{1}{2}||u_t - u_{t-1}||^2_{L_t}\right) \qquad (3.15)$$

the Maximum a Posteriori (MAP) estimate of $U_{1:T}$ is given to be:

$$\tilde{u}^*_{1:t} = \arg\max_{u_{1:t}} \pi_n(u_{1:t}) \qquad (3.16)$$

which essentially specifies the model that best explains the reference $y_{1:t}^{ref}$ as the output sequence from a general set of models defined in (3.2).

In [13], it had been shown that assuming a Markov transition density for U_t in (3.15), the optimal control defined in (3.5) and the MAP estimate defined in (3.16) are the same. Moreover, following the monotonicity of the transformation $p(\cdot)^\gamma$, the mode of $p(\cdot)^\gamma$ is the same for all $\gamma > 0$. The main difference is that it is easier to sample closer to the mode when $\gamma > 1$ because the density has sharper peak and vice versa. Putting all these together, we have the follows:

$$u^*_{1:t} = \tilde{u}^*_{1:t} = \arg\max_{u_{1:t}} \pi_n(u_{1:t})^\gamma, \ \gamma > 0 \qquad (3.17)$$

In SMC algorithm, this can be easily estimated as follows:

$$\hat{u}^*_{1:t} = arg\max_{u_{1:t}} \pi_n(u_{1:t}^{(i)})^\gamma = \arg\max_{u_{1:t}} p(u_{1:t}^{(i)} \mid y_{1:t}^{ref})^\gamma \qquad (3.18)$$

Algorithm 6 Rao-Blackwellised SMC to search for the optimal control parameters

---

1: **function** OPTIMALPARAMETERS(N, T)

2:    Set $t \leftarrow 1$.

3:    For $i \in 1, \ldots, N$, sample $u_1^{(i)} \sim q(u_1^{(i)} \mid y_1^{(i)})$.

4:    For $i \in 1, \ldots, N$, calculate the unnormalised importance weight:

$$\tilde{w}_1^{(i)} = \frac{p(u_1^{(i)})g_1(y_1 \mid u_1^{(i)})}{q_1(u_1^{(i)}, y_1)}$$

5:    For $i \in 1, \ldots, N$, normalize the importance weight:

$$\hat{w}_1^{(i)} = \frac{\tilde{w}_1^{(i)}}{\sum_{i=1}^{N} \tilde{w}_1^{(i)}}$$

6:    Set $t \leftarrow t + 1$.

7:    **while** $t \leq T$ **do**

8:        For $i \in 1, \ldots, N$, sample $u_t^{(i)} \sim q(u_t^{(i)} \mid y_{t-1}^{(i)}, u_{t-1}^{(i)})$.

9:        For $i \in 1, \ldots, N$, calculate the unnormalised importance weight:

$$\tilde{w}_t^{(i)} = w_{t-1}^{(i)} \frac{p_t(u_t^{(i)} \mid u_{t-1}^{(i)})g_t(y_t \mid u_{1:t}^{(i)}, y_{1:t-1})}{q_t(u_t^{(i)} \mid u_{1:t-1}^{(i)}, y_{1:t})}$$

   where

$$p(y_t \mid u_{1:t}^{(i)}, y_{1:t-1}) \sim N(m_{t|t-1}, S_t) \tag{3.12}$$

   with $m_{t|t-1}$ and S_t are updated with Kalman Filter recursions:

$$\mu_{t|t-1} = A_t(u_t)(\mu_{t-1|t-1})X_{t-1} + F_t(u_t)$$
$$\Sigma_{t|t-1} = A_t(u_t)\Sigma_{t-1|t-1}A_t(u_t)^T + B_t(u_t)B_t(u_t)^T$$
$$S_t = C_t(u_t)\Sigma_{t|t-1}C_t(u_t)^T + D_t(u_t)D_t(u_t)^T$$
$$m_{t|t-1} = C_t(u_t)\mu_{t|t-1} + G_t(u_t)$$
$$\mu_{t|t} = \mu_{t|t-1} + \Sigma_{t|t-1}C_t(u_t)S_t^{-1}(y_t - m_{t|t-1})$$
$$\Sigma_{t|t} = \Sigma_{t|t-1} - \Sigma_{t|t-1}C_t(u_t)S_t^{-1}C_t(u_t)\Sigma_{t|t-1}$$

10:        For $i \in 1, \ldots, N$, normalize the importance weight:

$$\hat{w}_t^{(i)} = \frac{\tilde{w}_t^{(i)}}{\sum_{i=1}^{N} \tilde{w}_t^{(i)}}$$

11:        **Resample:** For $i \in 1, \ldots, N$, resample $u_{1:t}^{(i)} \sim \frac{\sum_{i=1}^{N} \hat{w}_t^{(i)} \delta_{u_{1:t}^{(i)}}}{\sum_{j=1}^{N} \hat{w}_t^{(j)}}$.

12:        **Move:** For $i \in 1, \ldots, N$, sample $u_{1:t}^{(i)} \sim K_t(\cdot)$, where K_t is p_t-invariant.

13:    **end while**

14:    Compute the MAP estimate: $\hat{u}_{1:t}^* = \arg\max_{u_{1:t}} p(u_{1:t}^{(i)} \mid y_{1:t}^{ref})^\gamma$.

15: **end function**

As long as the support of the proposal density include the support of $(u_{1:t}^{(i)} \mid y_{1:t}^{ref})$, this estimate converge asymptotically to the $\tilde{u}_{1:t}^*$ as $N \to \infty$. A better estimate can be obtained by embedding the Viterbi strategy as discussed in [9], but this comes with a computational cost that is quadratic in terms of the number of samples.

To add diversity to the population samples, the Resample-Move step (see Section 2.5.6) can be added. This step perturbs the samples yet leave the distribution the samples represent unchanged by using a MCMC kernel that is invariant in density.

3.3   Numerical example: Tracking an oscillating wave

We first consider here a simple linear Gaussian state space model as follows:

$$X_t = X_{t-1} + W_t + U_t$$
$$Y_t = X_t + V_t \tag{3.19}$$

with $W_t, V_t \sim \mathcal{N}(0, I)$. This model is essential an instance of the class of model presented earlier in (3.2), with $A_t = B_t = C_t = D_t = I$, $F_t(U_t) = U_t$ and $G_t(U_t) = 0$. The target reference is set to be an oscillating wave: $y_t^{ref} = \cos(0.2\pi t + 0.3)$ and L_1 is set to 0.1 similar to the setting used in [13].

This model serves two purposes here. Firstly, it is sufficient simple to verify the implementation[1]. Secondly, it serves as good benchmark problem in which we can examine the sensitivity of the algorithm to different parameter settings.

We proceed by examining the proposed algorithm with proposal density $q_t(\cdot)$ set to be the same as the Markov transition density and *all* different combination of settings as follows:

1. Various time period length settings, T: 5, 10 and 20.

2. Various sample size settings, N: 100, 500, 1000, 5000 and 10000.

3. Resampling step: Always vs. Selectively (trigger threshold set to $\frac{\text{ESS}}{2}$).

4. Resample-Move step with MCMC (random walk proposal): Always vs. disabled.

5. Various γ settings: Constant function of 1, 50, 100, 1000 and increasing function of time t, $10t$, $50t$ and $100t$.

For each of the setting, the experiment is repeated for 30 times. Therefore, there are in total $3 \times 5 \times 2 \times 2 \times 8 \times 30 = 14400$ runs.

---

[1]Strictly speaking, testing only increases confidence of the implementation correctness but does not prove no bug.

3.3.1 Results and discussion

This sections presents a summary of the experimental results from all the runs. Instead of presenting the outputs produced by all the runs, only a subset of outputs that demonstrates our findings is presented here. The complete set of implementation code and output produced can be obtained from the repository of this project at GitHub.

3.3.1.1 Selectively resampling step with ESS

We first evaluate the effect of selectively enabling resample step with ESS mechanism. In *all* the settings, it is found that the simple runs with resample step enabled always at each iteration, i.e., with ESS disabled, have relatively better performance. Figure 3.1 shows the result the performance obtained in terms of $\log \pi_T(\hat{u}_{1:t}^*)$ of 30 independent runs of the algorithms with $T = 20$ and $N = 10000$ for various γ settings. All other runs with different settings of T and N concur with this finding. From here onwards, we assumes ESS mechanism is disabled unless stated otherwise.

To explain this counter-intuitive result, recall that ESS is essentially a proxy to measure the quality of the weighted samples in representing the target distribution. A sample set that is well spread over the target distribution would have high ESS value and vice versa. However, the application of SMC as a maximiser here is *only* concerned with the mode of the distribution. The process of resampling the samples proportional to the weights essentially pushes samples towards the mode of the distribution. Selectively resampling using ESS essentially reduces the strength of this push and therefore leads to sub-optimal solution.

3.3.1.2 Resample-Move step with MCMC

In Figure 3.1, we can also observe that the Resample-Move step does improve performance, albeit the improvement is marginal in most of the cases. The improvement is more significant when γ is set to be high. Figures 3.2, 3.3 and 3.4 show the box plots of $\log \pi_T(\hat{u}_{1:t}^*)$ of 30 independent runs for time period T of 5, 10, 20 respectively. Each sub-figure shows the runs for a selected γ setting with the box-plots grouped by whether the Resample-Move step is used (on the left) or not (on the right), sorted in terms of number of samples used N.

One possible explanation to this result is that high γ settings sharpen the distribution shape. Such a setting encourages optimisation by pushing the sample towards high density area but at the same time reduces the space exploration. With excessively high γ settings, the sample set may be led to a local optimal with relatively high density
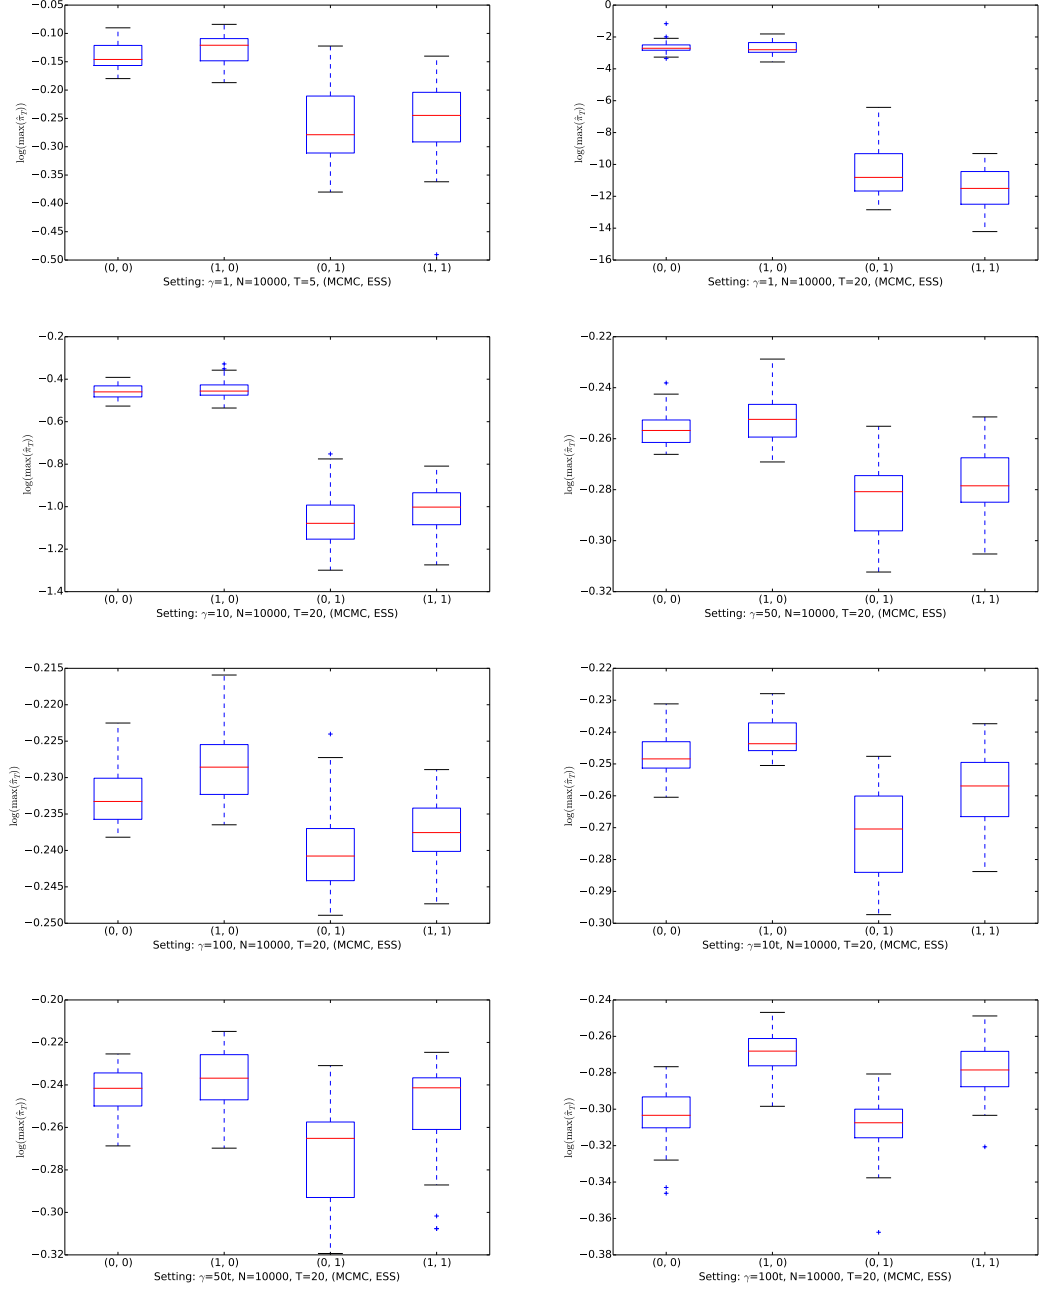
Figure 3.1: The performance in terms of $\log \pi_T(\hat{u}^*_{1:t})$ of 30 independent runs of the algorithms with $T = 20$ and $N = 10000$ for various γ settings. The four box-plots in each sub-figure represent the results of runs with the following settings (from left to right): a) MCMC disabled, ESS disabled, b) MCMC enabled, ESS disabled, c) MCMC disabled, ESS enabled, and d) MCMC enabled, ESS enabled.

area, which may not be the mode of the distribution. The perturbation Resample-Move step introduced at each iteration may help particles to escape from local optimal.
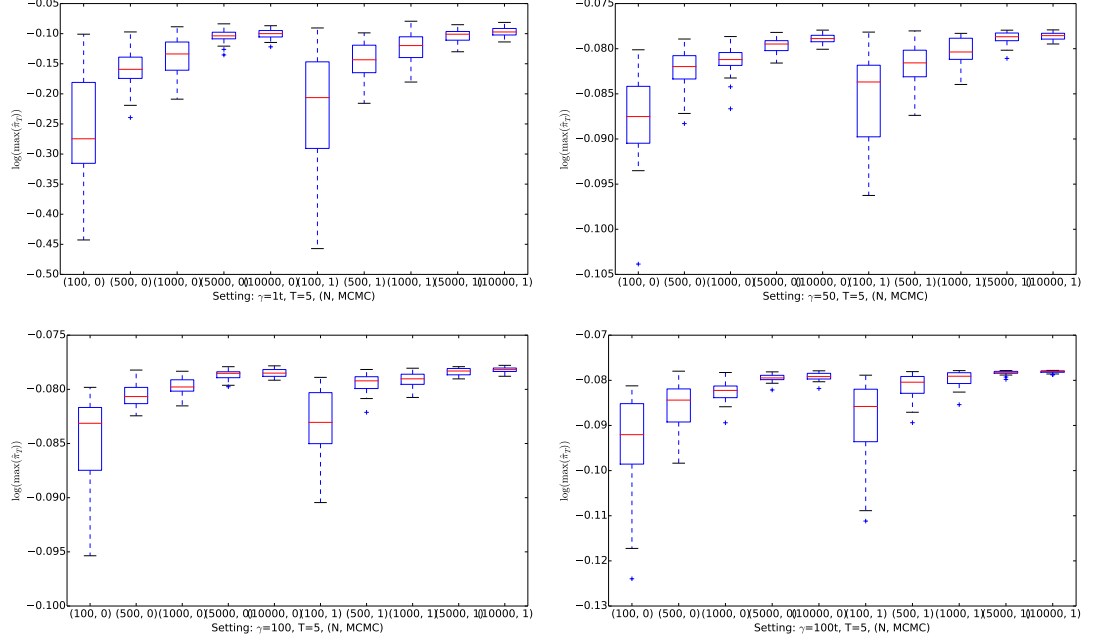

Figure 3.2: The $\log \pi_T(\hat{u}_{1:t}^*)$ box-plots of 30 independent runs with various settings for $T = 5$. Each sub-figure shows the runs for a selected γ setting with the box-plots grouped by whether the Resample-Move step is used (on the left) or not (on the right), sorted in terms of number of samples used N.

3.3.1.3 Various sample size settings

In terms of the sample size setting used, the performance increased monotonically with the number of samples used for all time periods considered as shown in Figures 3.2, 3.3 and 3.4. This is in-line with the theory, more samples is always better. For the shorter time-period considered, says $T = 5$, the performance improvement sometimes is "saturated" even with lower number of samples, says $T = 1000$ or $T = 5000$. This suggests the optimal control parameters may have been attained.

3.3.1.4 Various γ settings

A more interesting result is to compare the performance with different γ settings. Let first consider the fixed γ settings. Excessive high γ settings may lead to issue discussed earlier and stuck at local optimal. On the other hand, low γ settings also does not perform too well either (consider the process of searching for the mode by randomly
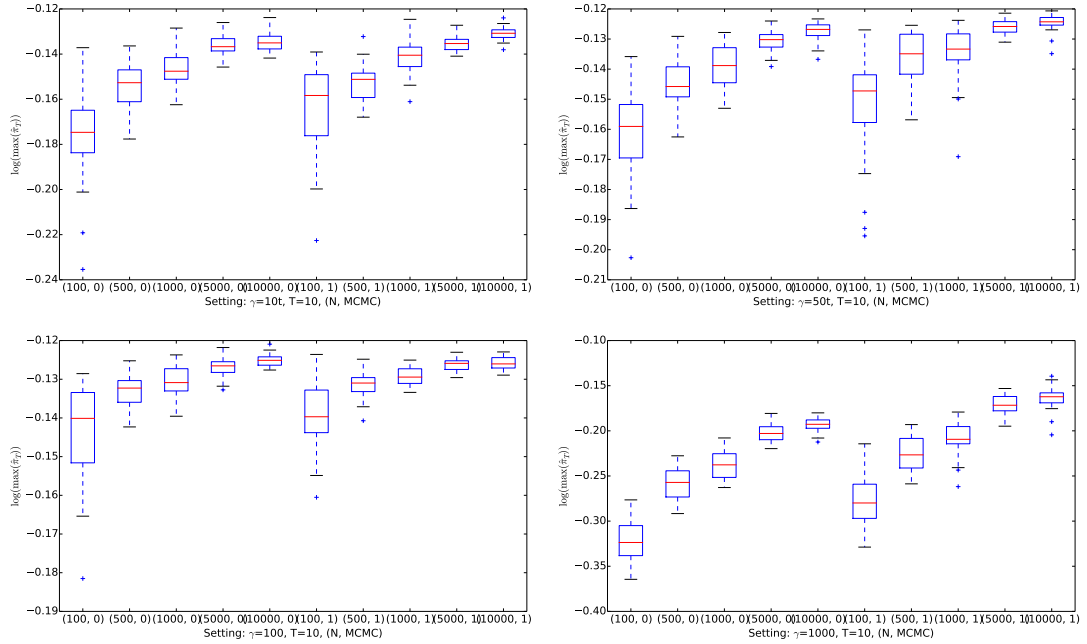
Figure 3.3: The $\log \pi_T(\hat{u}_{1:t}^*)$ box-plots of 30 independent runs with various settings for $T = 10$. Each sub-figure shows the runs for a selected γ setting with the box-plots grouped by whether the Resample-Move step is used (on the left) or not (on the right), sorted in terms of number of samples used N.

sampling a very flat distribution). The optimal γ setting seems to lie sometimes in between. In this case, the $\gamma = 100$ setting results in the best performance.

In [13], it was suggested that using an increasing γ setting may be a good pragmatic compromise between accuracy and good mixing. We investigate further on this by setting γ to be increasing function of the form αt, where α is constant and t is the time step. The results show that the performance seems improve when the α value is low, but the result is worse when the α value is high. This suggests a capping on the maximum value on α or a more moderate in t may be more suitable. We therefore carry out two extra γ settings: $50 \log(t + 1)$ and $100 \log(t + 1)$ to investigate this. The results of these runs are shown in Figure 3.5. The results obtained are not statistically different from a simple fixed $\gamma = 100$ setting.

Lastly, we conclude this section by showing the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in various runs in Figure 3.6 and their corresponding tracking errors in Figure 3.7. The results looks very promising. Except in the runs with the $\gamma = 1$ setting in which there is a lot of noise in tracking, all other runs are able to closely track the reference signal.
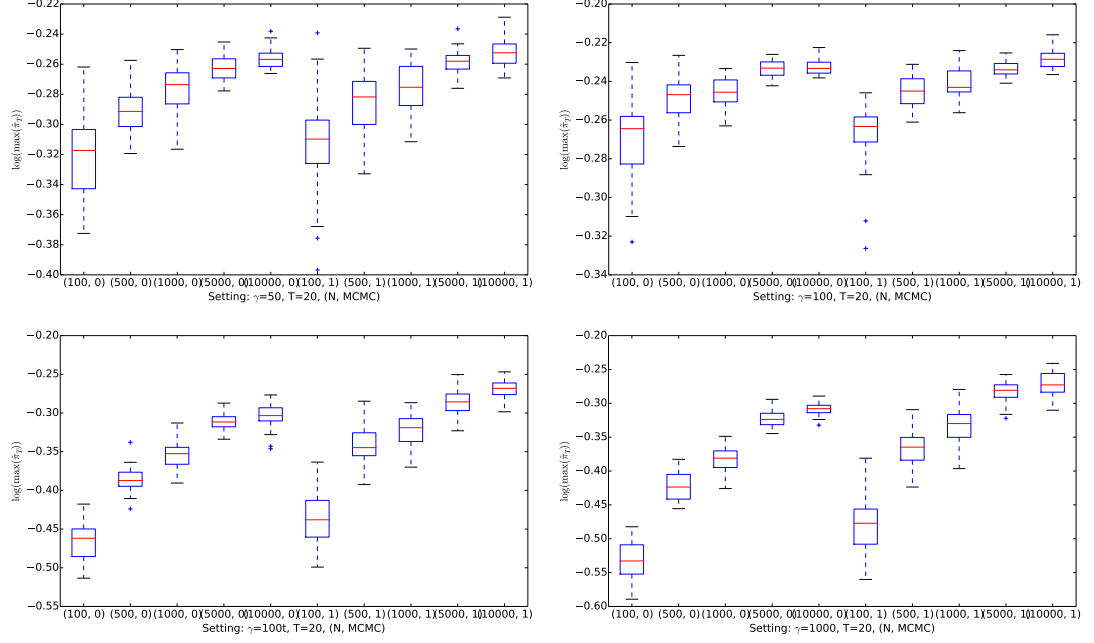
Figure 3.4: The $\log \pi_T(\hat{u}_{1:t}^*)$ box-plots of 30 independent runs with various settings for $T = 20$. Each sub-figure shows the runs for a selected γ setting with the box-plots grouped by whether the Resample-Move step is used (on the left) or not (on the right), sorted in terms of number of samples used N.

Figure 3.5: The $\log \pi_T(\hat{u}_{1:t}^*)$ box-plots of 30 independent runs with γ settings: $50 \log(t + 1)$ and $100 \log(t + 1)$. Each sub-figure shows the runs for a selected γ setting with the box-plots grouped by whether the Resample-Move step is used (on the left) or not (on the right), sorted in terms of number of samples used N.

29
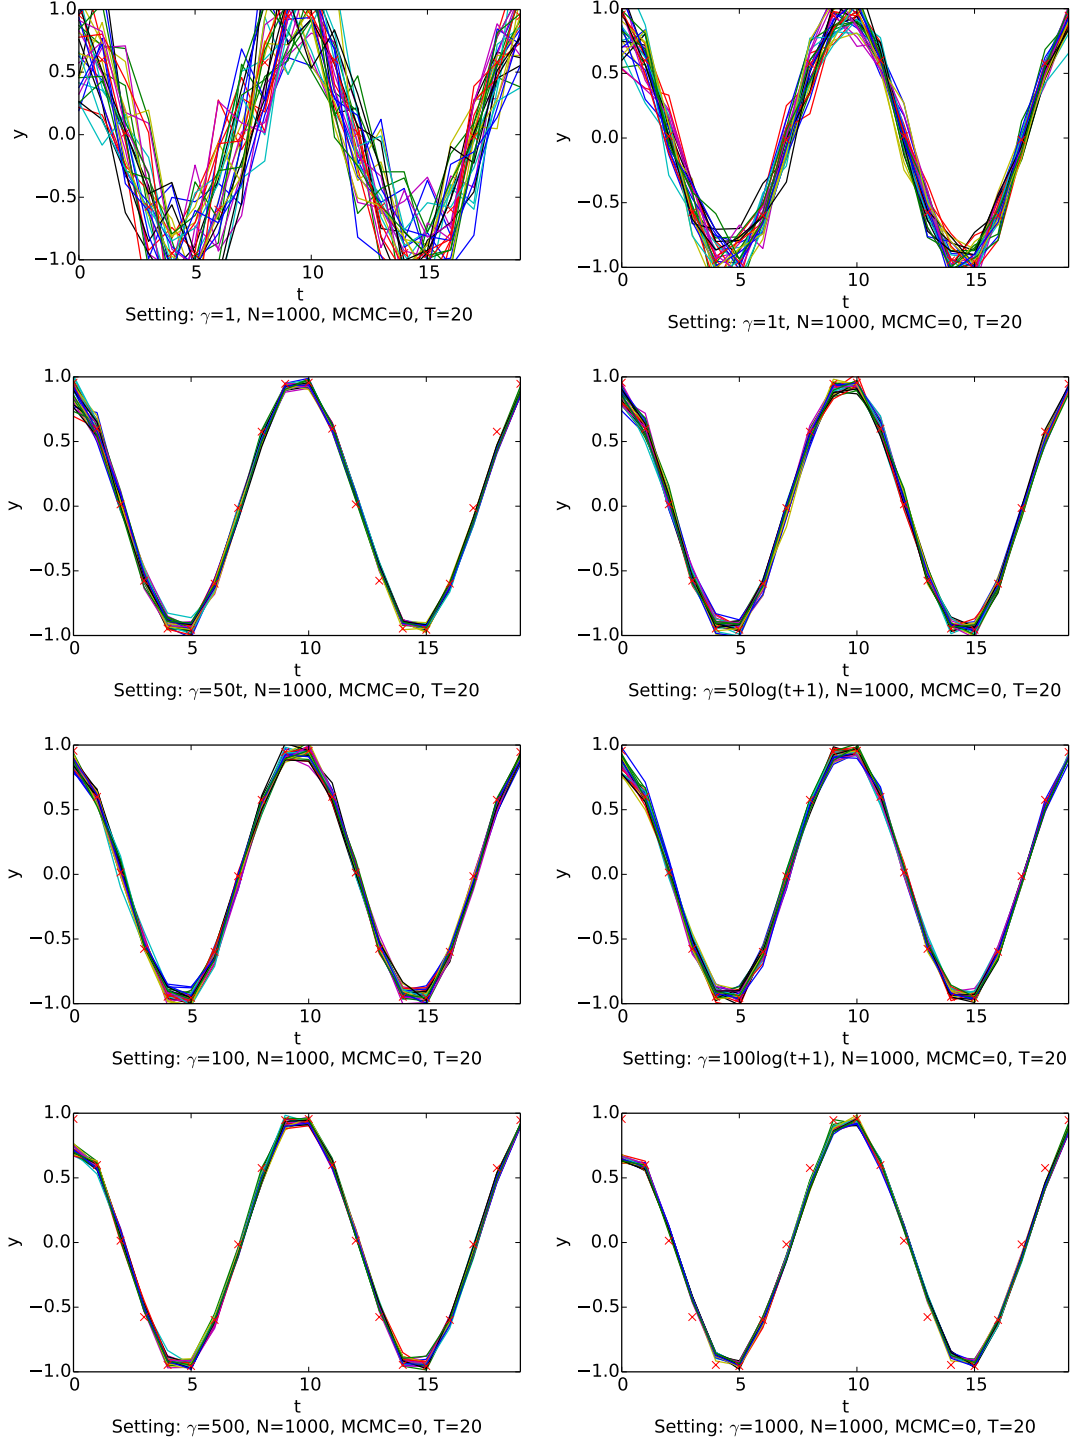
Figure 3.6: The means of the recursive likelihood $m_{t|t-1}(u^*_{1:T})$ obtained in 30 independent runs with $T = 20$, $N = 1000$, various γ settings and Resample-Move step disabled are plotted in lines with different colours. The target reference signal y^{ref} is also plotted with a 'x' marker.

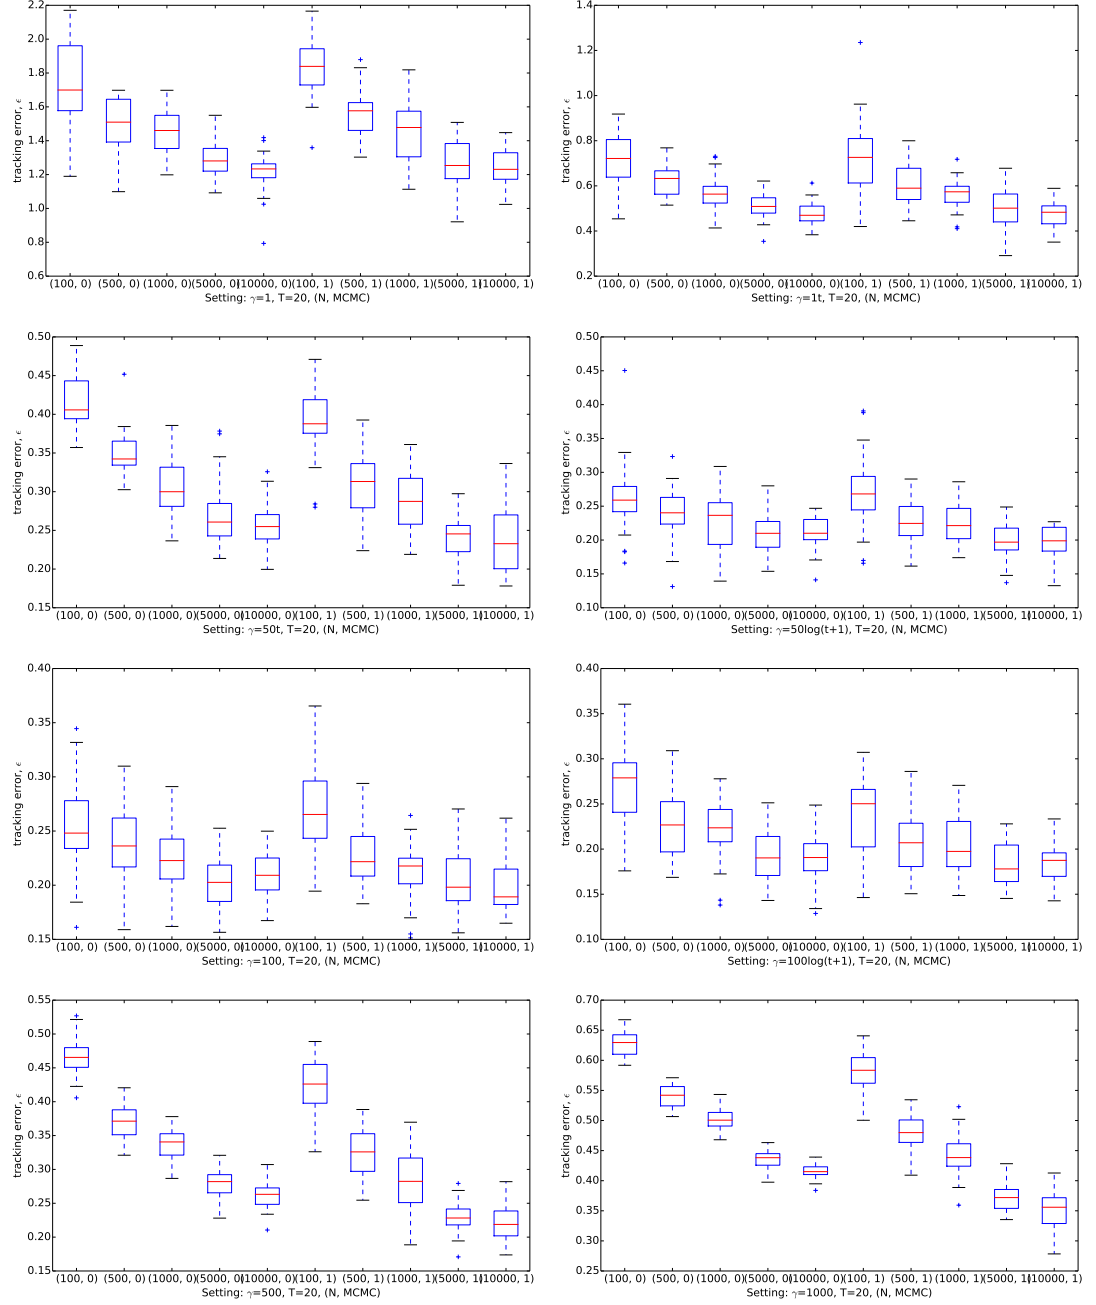Figure 3.7: The box-plots of tracking errors of the model using $u^*_{1:T}$ obtained in 30 independent runs with $T = 20$, $N = 1000$, various γ settings and Resample-Move step disabled.

3.4 Conclusions

In this chapter, we introduces the index tracking fund and summarise some causing factors of the tracking error between the fund and its benchmark index. We then present how the portfolio optimisation problem in terms of minimising the tracking error can be formulated as a stochastic control problem. This stochastic control problem can be turned into a path-space parameter estimation problem, in which Rao-Blackwellised SMC can be used to estimate these parameters efficiently. Lastly, we presents a some proof-of-concept experiment with a oscillating wave as target reference signal. The results show that the algorithm is able to track the reference signal well.

Chapter 4

Tracking the DAX index

In this chapter, we attempt to use the same algorithm to a real-world financial index. We look at the German's DAX (Deutscher Aktienindex) Index, which consists of 30 German's blue chip stocks listed on Frankfurt Stock Exchange as the constituents from 1st January 2014 up to 30th June 2014. It represents 80% of the aggregated prime standard's market capitalisation. The DAX Index is chosen for various pragmatic reasons. It is one of the major world indices that is tracked by various index funds, it has small number of components and the data is freely accessible[1]. For further detail on the DAX Index methodology, refer [1].

4.1 Tracking the DAX4 sub-index

In the first experiment, we define a simple hypothetical sub-index, DAX4 consists of four stocks from the DAX Index that have the highest weights as of 2nd January 2014, namely Bayer (BAYN), Siemens (SIE), BASF (BAS) and Daimler (DAI). Then, the DAX4 index level is calculated as the simple weighted average of the adjusted close prices obtained from Yahoo Finance as follows:

$$y^{DAX4} = \sum_{s \in \mathbb{S}} w_s x_s \tag{4.1}$$

where y^{DAX4} is the index level, $\mathbb{S}$ consists of the four stocks and w_s and x_s are the weight and price of stock s respectively. Here the weights of these four stocks are assigned to be 0.4, 0.3, 0.2 and 0.1 respectively. The adjusted close price of each stock along with the calculated index level are shown in Figure 4.1.

---

[1]Actually, we first looked at Dow Jones Industrial Average (DJIA) Index obtained. It was later found some of the data is not publicly available. Having said that, the preliminary findings concur with the findings we have with the DAX Index reported here.

Figure 4.1: The adjusted close price of the 4 stocks and the calculated DAX4 index level, y^{DAX4}.

The portfolio optimisation problem is formulated as such y^{DAX4} is viewed as the target reference level that a portfolio manager attempt to replicate as close as possible by changing the holding position on each component stocks in $\mathbb{S}$, at the same time, minimise the transaction cost incurred in position changes.

To solve this problem using the SMC, the following state space modelling is used:

$$X_t = X_{t-1} + F_t(U_t) + W_t \tag{4.2}$$

$$Y_t = U_t^T X_t + 0.01 V_t \tag{4.3}$$

where $W_t \sim \mathcal{N}(\mu_{t_0}, \Sigma_{t_0})$, $V_t \sim \mathcal{N}(0, I)$, $\{X_t\}_{t\geq 0}$ is a vector of stock price processes modelled as Arithmetic Brownian Motion with drift, $\{U_t\}_{t\geq 0}$ is a vector of control input processes, each represents the position we have for each stock, $F_t(U_t)$ can be viewed as the market impact on price due to position changes and is set to be $0.0001U_t$ here, μ_{t_0} and Σ_{t_0} are vector of the estimated mean returns and the estimated covariance matrix of the stock returns and $\{Y_t\}_{t\geq 0}$ is the process represents the index level.

The values of μ_t and Σ_t here are estimated using Exponential Weighted Moving Average (EWMA) approach with the decay rate, λ, set to 0.94. For example, μ_t can

be easily calculated in a recursion form as follows:

$$\mu_t = \lambda \mu_{t-1} + (1 - \lambda) r_t \tag{4.4}$$

where r_t is the return at time step t. It is obvious from these equations that the EWMA estimate at time t depends on all the preceding estimates at time s, where $s < t$. To ensure the quality of EWMA estimate, a 6 months warm up period is used, i.e., the EWMA estimate is calculated from 1st July 2013 onwards. The estimates of μ_{t_0} and Σ_{t_0} obtained are as follows:

$$
\mu_{t_0} = \begin{pmatrix} 0.03125069 \\ 0.20020445 \\ 0.08702132 \\ 0.18319923 \end{pmatrix}
\Sigma_{t_0} = \begin{pmatrix}
0.60011676 & 0.69688309 & 0.44198994 & 0.66709016 \\
0.69688309 & 1.34168488 & 0.58394515 & 0.76533660 \\
0.44198994 & 0.58394515 & 0.48855826 & 0.56500417 \\
0.66709016 & 0.76533660 & 0.56500417 & 1.13794411
\end{pmatrix}
\begin{matrix} BAYN \\ SIE \\ BAS \\ DAI \end{matrix}
$$

(with column headers $BAYN$, SIE, BAS, DAI)

$$\tag{4.5}$$

It is worth here to re-iterate here that this model is just a means to an end to demonstrate the idea. Other sophisticated models, e.g., Geometric Brownian Motion with drift, Jump diffusion model, etc., are possible, perhaps better.

Using this model, we can write the reward function as follows:

$$J(u_{1:T}, y_{1:T}^{ref}, x_0) = \mathbb{E}_{x_0} \left[\exp\left(-\frac{1}{2} \sum_{t=1}^{T} \left(||y_t^{ref} - u^T x_t||_{Q_t(u_t)^{-1}}^2 + ||u_t - u_{t-1}||_{L_t}^2 \right) \right) \right] \tag{4.6}$$

Here, we define $L = 0.1 \operatorname{diag}(4, 3, 2, 1)$ and also impose additional constraints on the holding position of each individual component of U_t such that the range is within $[0, 0.5]$. The idea behind these choices is to construct a scenario in which there is different amount of transaction cost involved in changing the holding position on each individual component and a constraint on the maximum amount on the holding position of each individual component one portfolio can have.

For simplicity, we will use the SMC algorithm to target π directly, i.e, fixed $\gamma = 100$ setting and important sampling step, $U_{1,n}$ is sampled uniformly from the range bounded the constraints mentioned. The experiment is carried out for 30 independent runs, each with sample size is fixed to 10000 and no Resample-Move step is used here.

4.1.0.5 Results and discussion

As in previous experiment, the estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs are shown in Figure 4.4. It is obvious the output of the model produce using the estimated $\hat{u}_{1:T}^*$

Figure 4.2: The estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs with the y^{DAX4} set to be the target reference using $L = 0.1 \operatorname{diag}(4, 3, 2, 1)$.

is closely tracking the reference index level, y^{DAX4}. The average annualised[1] tracking error of 0.023288 over the 30 runs. Having said that, the estimated set of control parameter $\hat{u}_{1:T}^*$ which represents the holding position of each stock does not seem very stable.

To reduce the number of changes in holding positions, the value of L matrix can be scaled up. To demonstrate this, the same experiment is re-run with $L = 5 \operatorname{diag}(4, 3, 2, 1)$. The results are shown in Figure 4.3. It is obvious the estimated set of control parameters $\hat{u}_{!:T}^*$ produced in this case is much more stable. However, this comes at a price on the tracking error. The average annualised tracking error with this setting is increased to 0.036386 over the 30 runs.


Figure 4.3: The estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs with the y^{DAX4} set to be the target reference using $L = 5 \operatorname{diag}(4, 3, 2, 1)$.

---

[1] the annualisation factor is set to be $\sqrt{252}$ as the daily return is measured here.

4.1.1 Tracking the DAX index

Given the results of the experiment for DAX4 index look promising, we now attempt to track the original DAX index which consists of 30 components using a similar model as follows:

$$X_t = X_{t-1} + F_t(U_t) + W_t \tag{4.7}$$

$$Y_t = 30U_t^T X_t + 0.01V_t \tag{4.8}$$

where $W_t \sim \mathcal{N}(\mu_{t_0}, \Sigma_{t_0})$, $V_t \sim \mathcal{N}(0, I)$, along with the same parameter settings as before except the following:

1. The target reference y^{ref} is the daily close level of the DAX index level.

2. The number of dimensions for U_t and X_t are now 30.

3. The L matrix is set to be $5I_{30 \times 30}$, where $I_{n \times n}$ is the n dimensional identity matrix, i.e., the transaction cost of each component is equal.

As mentioned in Section 3.1.1, sometimes it can be more efficient not only just in terms of computational, but also in terms of transaction cost and management perspective to track the benchmark index with only a subset of its component stocks. To do this, only minimal changes are required to apply on the model proposed. We choose to use the top 6 highest weighted constituents of DAX index by including Allianz SE (ALV) and SAP SE (SAP), in which they all together constitute more than 50% of the weights for the DAX index level as of 2nd January 2014. Using only these 6 component stocks and a slight increase on the range of each component of U_t to within $[0, 1]$, the experiment is re-run for 30 times.

4.1.1.1 Results and discussion

As before, the estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs with the DAX index level set to be the target reference for both the full replication setting and partial replication setting are shown in Figure 4.4 and 4.5 respectively.

In the full replication setting, the average annualised tracking error over 30 runs is found to be 0.113118. However, the computational cost increases significantly in this experiment due to the dimensionality increases in U and X. Using the computation cost of the previous $DAX4$ experiment as the baseline, the computational cost for this experiment is ≈ 7.2 times more per time step.

In the partial replication setting, the average annualised tracking error over 30 runs is found to be 0.062650. The estimated optimal control $u_{1:T}^*$ found in this setting is also much stable. This means smaller changes in holding positions and therefore lesser transaction costs. In terms of computational cost, the relative computational cost factor in comparison to the baseline is approximately 1.1. This translates to approximately $6.5\times$ speed-up in comparison to the full replication setting.


Figure 4.4: The estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs with the DAX index level set to be the target reference using all 30 component stocks.


Figure 4.5: The estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs with the DAX index level set to be the target reference using only the 6 component stocks with highest weights.

4.2 Model Predictive Control (MPC)

Up to this point, the discussion has been focussed on predicting the optimal (as outlined in the reward function) strategy for the a given time period $0:T$. However, as time progress, more realisation of the reference signal $y_{t'}^{ref}, t' < T$ are observed. It makes

sense to take this information into account in the optimisation process. In this section, we shall demonstrate how the SMC technique can be integrated into Model Predictive Control (MPC) framework.

MPC is an established technique in control theory which has achieved significant success in many different domains. It is a kind of receding horizon control algorithm, in which at time step t, the model is used to optimise the controls for the process, $u_{t:t+k}$ for a finite horizon for $t : t + k$. However, only the first set of controls, i.e., u_t is applied to the system. At time $t + 1$, the new set of the observation is measured to update the states. Then, the whole optimisation process is repeated for the same horizon, but shifted by one time step until $t = T$. The MPC framework is summarised in Algorithm 7. Refer [16, 19] for further details on MPC.

Algorithm 7 Model Predictive Control

1: **function** MODELPREDICTIVECONTROL(T,k)
2:    Set $t \leftarrow 1$.
3:    **while** $t \leq T$ **do**
4:        Search the optimal $u^*_{t:t+k}$ for the control problem of time horizon $t : t + k$.
5:        Apply the first set of the optimal control u^*_t to the problem.
6:        Update the model states with any new measurement y^{ref}_t.
7:        Set $t \leftarrow t + 1$.
8:    **end while**
9: **end function**

To use MPC framework here, the SMC based optimisation algorithm can be executed on a daily basis after the market close on day t to search for the optimal $u^*_{t:t+k}$, but only apply the first set of controls, u_t to the problem. The same algorithm is repeated on the next day with a shifted time horizon and so on. Using this model, the reward function that we attempt to optimise at each time step t is as follows:

$$J(u_{t:t+k}, y^{ref}_{t:t+k}, x_t) = \mathbb{E}_{x_0} \left[\exp \left(-\frac{1}{2} \sum_{s=t}^{t+k} \left(||y^{ref}_s - u^T x_s||^2_{Q_s(u_s)^{-1}} + ||u_s - u_{s-1}||^2_{L_s} \right) \right) \right]$$
$$(4.9)$$

To demonstrate this idea, we re-run the DAX index with partial replication setting using this MPC framework. The setting of k is set to be 5. The experiment is repeated for 30 times. The results look very promising. Figure 4.6 shows the estimated optimal control $u^*_{1:T}$ and the mean of the recursive likelihood $m_{t|t-1}(u^*_{1:T})$ obtained in one of the independent runs. The tracking performance looks good, with an average annualised tracking error of 0.062813 over the 30 runs. The estimated optimal control $u^*_{1:T}$ found

is also reasonable stable. In terms of computational cost, the relative computational cost factor in comparison to the baseline is approximately 5.5 for each time step.


Figure 4.6: The estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs with the DAX index level set to be the target reference using the MPC framework with $k = 5$.

In practice, the target reference signal $y_{t+1:t+k}^{ref}$ would not be observable at time-step t. To over this, we first define an estimator of $y_{t+1:t+k}^{ref}$ to be as follows:

$$\hat{y}_{t+1}^{ref} = y_t^{ref} + \beta k \tag{4.10}$$

where β is the EWMA estimate ($\lambda = 0.94$) of the daily changes of y^{ref} up to time step t. Again, this estimator is used just for simplicity sake. Other estimators are possible.

Using the estimate $\hat{y}_t^{ref}$ in place of y_t^{ref}, the reward function J defined in (4.9) now only depends information up to time step t. With this setting, the experiment is re-run for 30 times. Figure 4.7 shows the estimated optimal control $u_{1:T}^*$ and the mean of the recursive likelihood $m_{t|t-1}(u_{1:T}^*)$ obtained in one of the independent runs. The tracking performance looks good, with an average annualised tracking error of 0.062364 over the 30 runs. The estimated optimal control $u_{1:T}^*$ found is also reasonable stable.

4.3 Conclusions

This chapter presents several experiments in apply the SMC technique to track a real-world financial index — the DAX index. It begins with a simple experiment in which a synthetic sub-index constructed with only 4 to proof the concept. It then uses the same technique to demonstrate how the DAX index can be tracked with full replication as well as partial replication. Lastly, it introduces Model Predictive Control (MPC) framework and demonstrates how this technique can be integrated into the MPC framework easily. In all cases, the results sow the SMC technique proposed is able to search for optimal control that results an output that track the reference signal given.

Figure 4.7: The estimated optimal control $u^*_{1:T}$ and the mean of the recursive likelihood $m_{t|t-1}(u^*_{1:T})$ obtained in one of the independent runs with the estimated DAX index level $(\hat{y}^{ref}_{t:t+k})$ set to be the target reference using the MPC framework with $k = 5$.

Chapter 5

Conclusions and Future Work

This thesis demonstrates how SMC can be used as an effective means of searching the optimal strategy for index tracking funds in terms of tracking error and transaction cost. The sensitivity of the technique with respect to the parameter settings and annealing techniques are explored. Several numerical examples using real-world data have been presented, showing how an index can be fully replicated or partially replicated with transaction cost taken into account. It also presents the Model Predictive Control (MPC) framework and demonstrates how the SMC technique can be integrated into MPC framework. The experimental results show the technique is very promising.

There are numerous possible directions for future work that have been identified during the course of this research.

1. More realistic models — The Arithmetic Brownian Model with drift used in this thesis is rather simple. A possible extension work is to consider a more advanced model, e.g., Geometric Brownian Model, Jump Diffusion model, etc. Moving away from conditional Gaussian model is also possible. However, the inner Kalman Filter recursion is no longer optimal outsides the Gaussian assumption. A possible solution to this is substituting the Kalman Filter with a nested SMC algorithm. This setup is known as the SMC2 algorithm [5].

2. Parallel computation — The nested SMC setup inevitably adds consideration amount of computation requirements. A possible way to speed-up is to parallelise some of the steps in SMC algorithm. This is very straight-forward for many steps, except the resampling step, which remains an interesting research topic on its own.

3. More complex financial indices — The benchmark index used in this thesis is rather simple. This can be potentially an issue. Having said that, this index is a simple, but by no means a "toy" index. It is a major financial index in the world.

There are however indices with much large number of components, e.g., ≈ 1600 for MSCI World Index. This translates to a high dimensional problem in which may be difficult for the proposed model here. A possible way to cope with this is to just track the index with partial replication using pre-selected subset of components, perhaps using Principle Component Analysis (PCA). Alternatively, we could use a divide and conquer approach, tracking multiple sub-indices separately. Yet, there is still much to answer here, for example:

- How to split an index into smaller components in a systematic manner?

- How to deal with index components changes during announcement and implementation?

5.1 Closing remarks

The work reported in this thesis demonstrates a considerable degree of originality supported by extensive experimentation. The case studies are necessarily limited given the limited amount of time frame. However, the results demonstrate that portfolio optimisation using Sequential Monte Carlo techniques have very considerable promise. We recommend these approaches to the research community for further investigation.

Appendix A

Kalman Filter

In a conditional linear Gaussian state space as follows:

$$X_t^L = A_t(X_t^N)X_{t-1}^L + B_t(X_t^N)W_t + F_t(X_t^N) \tag{A.1}$$

$$Y_t = C_t(X_t^N)X_t^L + D_t(X_t^N)V_t + G_t(X_t^N) \tag{A.2}$$

where $\{X_t^N\}_{t\geq0}$ is a non-linear Markov process, A_t, B_t, C_t, D_t, F_t, G_t are appropriate matrix/vector of X_t^N and $\{W_t\}_{t\geq0}$ and $\{V_t\}_{t\geq0}$ are independent sequences of standard Gaussian random variable, i.e., $W_t, V_t \sim \mathcal{N}(0, I)$. In such a case, the transition density and likelihood of this model are Gaussian distributions with centre lied at a point of a linear combination of the known x_t^N of the following form:

$$f_t(x_t^L \mid x_{t-1}^L, x_t^N) = \mathcal{N}(A_t(x_t^N)x_{t-1}^L + F_t(x_t^N), B_t(x_t^N)B_t(x_t^N)^T)$$

$$g_t(y_t^L \mid x_t^L, x_t^N) = \mathcal{N}(C_t(x_t^N)x_t^L + G_t(x_t^N), D_t(x_t^N)D_t(x_t^N)^T) \tag{A.3}$$

Using the properties of Gaussian distribution, the integral involved can be resolved analytically. This leads to the widely used *Kalman Filter* [12] that has the following recursive solution as follows:

$$\mu_{t|t-1} = A_t(x_t^N)(\mu_{t-1|t-1})X_{t-1} + F_t(x_t^N) \tag{A.4}$$

$$\Sigma_{t|t-1} = A_t(x_t^N)\Sigma_{t-1|t-1}A_t(x_t^N)^T + B_t(x_t^N)B_t(x_t^N)^T \tag{A.5}$$

$$S_t = C_t(x_t^N)\Sigma_{t|t-1}C_t(x_t^N)^T + D_t(x_t^N)D_t(x_t^N)^T \tag{A.6}$$

$$m_{t|t-1} = C_t(x_t^N)\mu_{t|t-1} + G_t(x_t^N) \tag{A.7}$$

$$\mu_{t|t} = \mu_{t|t-1} + \Sigma_{t|t-1}C_t(x_t^N)S_t^{-1}(y_t - m_{t|t-1}) \tag{A.8}$$

$$\Sigma_{t|t} = \Sigma_{t|t-1} - \Sigma_{t|t-1}C_t(x_t^N)S_t^{-1}C_t(x_t^N)\Sigma_{t|t-1} \tag{A.9}$$

where $\mu_{t|t-1}$ and $\Sigma_{t|t-1}$ are the predicted mean and co-variance matrix of the state x_t^N, $m_{t|t-1}$ and S_t are the mean and co-variance matrix of the measurement at time t and

$\mu_{t|t}$ and $\Sigma_{t|t}$ are the estimated mean and co-variance matrix of the state x_t^N after seeing the observation y_t.

There are various extensions have been developed on top of this approach. For example, the Extended Kalman Filter (EKF) which uses Taylor Series expansion to linearise at the conditional variables locally [23], Unscented Kalman Filter which further extend the idea in EKF by only using a minimal set of well chosen samples [22], etc.

References

[1] Factsheet DAX. Websites, June 2014. 33

[2] ANDRIEU, C., DOUCET, A., AND HOLENSTEIN, R. Particle Markov chain Monte Carlo methods. *Journal of the Royal Statistical Society Series B 72*, 3 (2010), 269–342. 16

[3] BERZUINI, C., AND GILKS, W. RESAMPLE-MOVE Filtering with Cross-Model Jumps. In *Sequential Monte Carlo Methods in Practice*, A. Doucet, N. de Freitas, and N. Gordon, Eds., Statistics for Engineering and Information Science. Springer New York, 2001, pp. 117–138. 14

[4] CAPPÉ, O., MOULINES, E., AND RYDEN, T. *Inference in Hidden Markov Models (Springer Series in Statistics)*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2005. 8, 16

[5] CHOPIN, N., JACOB, P. E., AND PAPASPILIOPOULOS, O. SMC2: An efficient algorithm for sequential analysis of state space models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology) 75*, 3 (2013), 397–426. 42

[6] DOUC, R., AND CAPPE, O. Comparison of resampling schemes for particle filtering. In *Image and Signal Processing and Analysis, 2005. ISPA 2005. Proceedings of the 4th International Symposium on*. 13

[7] DOUCET, A., AND JOHANSEN, A. M. A tutorial on particle filtering and smoothing: Fifteen years later. *Handbook of Nonlinear Filtering 12* (2009), 656–704. 6

[8] GILKS, W. R., AND BERZUINI, C. Following a moving targetMonte Carlo inference for dynamic Bayesian models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology) 63*, 1 (2001), 127–146. 14

[9] GODSILL, S., DOUCET, A., AND WEST, M. Maximum a Posteriori Sequence Estimation Using Monte Carlo Particle Filters. *Annals of the Institute of Statistical Mathematics 53*, 1 (2001), 82–96. 24

[10] JOHANSEN, A. M., EVERS, L., AND WHITELEY, N. Monte Carlo Methods. *Lecture Notes 200* (2010). 13, 14, 15

[11] JOHNSON, B., BIOY, H., KELLETT, A., AND DAVIDSON, L. On the Right Track: Measuring Tracking Efficiency in ETFs. 18, 19

[12] KALMAN, R. E. A New Approach to Linear Filtering and Prediction Problems. *Transactions of the ASME–Journal of Basic Engineering 82*, Series D (1960), 35–45. 44

[13] KANTAS, N., MORAL, P. D., AND VINTER, R. Particle methods for Stochastic Regulation: towards an application for power scheduling. Tech. rep., Department of Electrical and Electronic Engineering, Imperial Colege, London, March 2011. 20, 22, 24, 28

[14] KITAGAWA, G. Monte Carlo Filter and Smoother for Non-Gaussian Nonlinear State Space Models. *Journal of Computational and Graphical Statistics 5*, 1 (1996), 1–25. 13

[15] LIU, J. S., AND CHEN, R. Sequential Monte Carlo Methods for Dynamic Systems. *Journal of the American Statistical Association 93* (1998), 1032–1044. 13

[16] MACIEJOWSKI, J. *Predictive Control: With Constraints.* Pearson Education. Prentice Hall, 2002. 39

[17] MARKOWITZ, H. Portfolio Selection. *The Journal of Finance 7*, 1 (1952), 77–91. 1, 17

[18] RABINER, L. R. A tutorial on hidden Markov models and selected applications in speech recognition. In *Proceedings of IEEE* (1989), pp. 257–286. 16

[19] RAWLINGS, J., AND MAYNE, D. *Model Predictive Control: Theory and Design.* Nob Hill Pub., 2009. 39

[20] ROBERT, C. P., AND CASELLA, G. *Monte Carlo Statistical Methods (Springer Texts in Statistics).* Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2005. 6

[21] ROCKAFELLAR, R. T., AND URYASEV, S. Optimization of Conditional Value-at-Risk. *Journal of Risk 2* (2000), 21–41. 1, 2

[22] WAN, E. A., AND MERWE, R. V. D. The Unscented Kalman Filter. In *Kalman Filtering and Neural Networks* (2001), Wiley, pp. 221–280. 45

[23] WELCH, G., AND BISHOP, G. An Introduction to the Kalman Filter. Tech. rep., University of North Carolina at Chapel Hill, Chapel Hill, NC, USA, 1995. 45

[24] WHITTLE, P. *Risk-sensitive Optimal Control*. Wiley-Interscience series in systems and optimization. Wiley, 1990. 20